

CS F433 Computational Neuroscience Assignment Report

Pratyush Bindal*, Kalash Bhattad*

Abstract- Hopfield Networks are associative memory models capable of storing and recalling information. In this study, we implement and evaluate a Hopfield Network trained on 16 x 16 binary images using the Hebbian Learning Rule^[1]. The network was tested under two distinct corruption methods - pixels randomly flipping with probability p and cropping images where only a central bounding box is preserved, and in contrast, surrounding pixels are set to a uniform colour. We investigate the network's ability to recover original patterns through synchronous and asynchronous update rules, analysing convergence behaviour and quantifying the fraction of correctly recovered images, providing insights into the stability and recall capacity of the Hopfield Network.

Keywords- Hopfield Networks, associative memory, Hebbian Learning Rule, synchronous updates, asynchronous updates

I. DATASET PREPARATION

We constructed the following dataset consisting of 25 16 X 16 pixels black and white images in the binary format. Furthermore, since PBM files consist of binary format (0/1), we converted them into a bipolar representation (-1/1).




Fig 1. Generated 16 x 16 Binary Patterns

Stored patterns influence Hopfield Network's stability and recall performance. Highly correlated patterns introduce interference, which may lead to spurious attractors and reduce retrieval accuracy.

Hence, we visualised the hamming distance and cosine similarity between the patterns to check for the level of correlation.

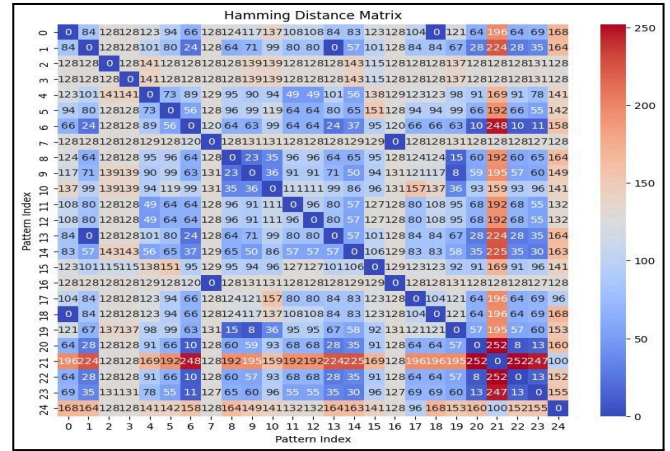


Fig 2. Hamming Distance Matrix

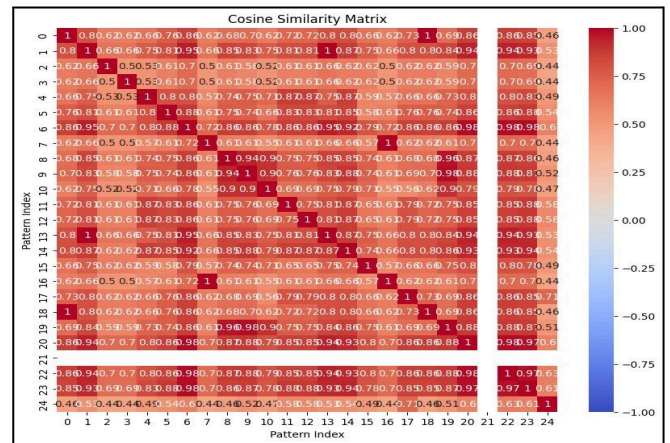


Fig 3. Cosine Similarity Matrix

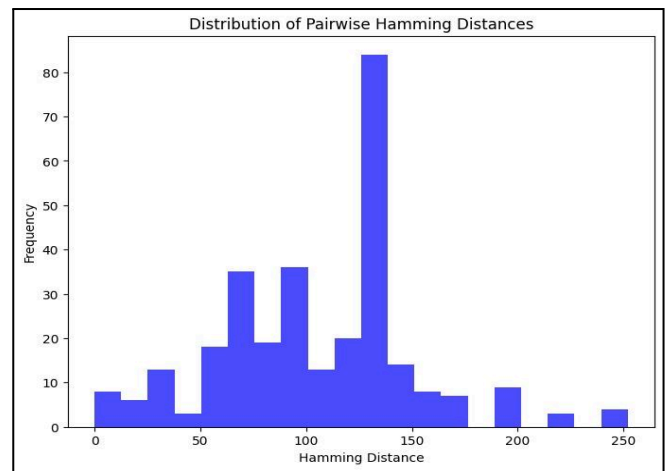


Fig 4. Distribution of pairwise Hamming Distances (Number of differing bits vs. Count of occurrences of each distance)

Furthermore, a fully connected Hopfield Network using the Hebbian Learning Rule with N neurons can theoretically store $N/(2 \ln N)$ patterns before retrieval performance degrades due to spurious attractors and retrieval errors^{[2][3]}. For a 16×16 binary image, the network consists of 256 neurons, which means the theoretical storage limit is around 16 patterns. The limit further degrades for practical purposes due to the correlation between images. Hence, we trained and tested the built Hopfield network for 25, 10 and 5 images out of all 25 images in the dataset for a comprehensive comparison and analysis.

II. IMPLEMENTATION AND SIMULATION OVERVIEW

A. Hopfield Network Training

We trained the Hopfield Network using the Hebbian Learning Rule^[3], which is local and incremental and updates the weight matrix W as follows: $W = \frac{1}{N} \sum \epsilon_i^u \epsilon_j^u$ where ϵ_i^u represents bit i from pattern u , and N is the number of training images. The diagonal elements of W were set to zero to avoid self-feedback.

B. Pattern Corruption

We used two corruption methods to test the network:

1. **Pixel Flipping:** Each pixel was flipped with a probability p where $p \in \{0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9\}$, achieved by randomly inverting pixel values in the images.
2. **Cropping:** A 10×10 bounding box was retained, while the outer pixels were set to black (-1) or white (+1).

C. Update Mechanisms and Convergence

We implemented two update methods on the network:

1. **Synchronous Update:** All neurons were updated simultaneously.
2. **Asynchronous Update:** We updated neurons one at a time in a random order, using the weighted sum of their connections to decide their new state.

The network was updated until convergence, that is, there are no further changes in state. Furthermore, intermediate states were saved as PBM files at the midpoint of the update process to analyse the modification of the Hopfield network.

D. Experimental Analysis

Experiments were conducted for $p \in \{0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9\}$, with 20 trials per p value as stated in the problem statement. We evaluated the results obtained based on two performance metrics:

- **Average convergence steps:** The number of updates required to reach a stable state.
- **Fraction of correctly recovered patterns:** The fraction of runs where the final state matched the original stored pattern.

III. RESULTS

We evaluated the performance of the model for 25, 10 and 5 sets of images. While testing on the entire dataset of 25 images, the Hopfield network performed suboptimally, verifying the theoretical limit for the Hopfield network as described^{[2][3]}. The following bar graph visualises the performance by showcasing average convergence steps reached across varying image bit corruption probabilities for both synchronous and asynchronous updates.




Fig 5. Average convergence Steps vs. Corruption probability for 25 image dataset

Additionally, we analysed the recall performance by analysing the fraction of correctly retrieved patterns across varying image bit corruption probabilities for both synchronous and asynchronous updates.

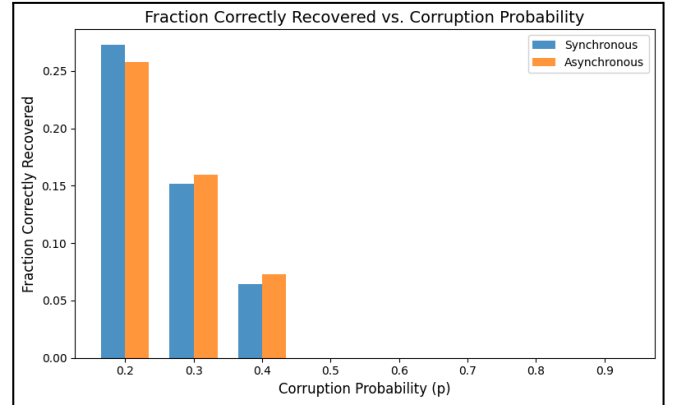


Fig 6. Fraction of correctly recovered images vs. Corruption probability for 25 Images Dataset

Implicitly, the fraction correctly recovered for an image dataset greater than the theoretical limit was highly erroneous, giving zero or almost zero correctness as the corruption was increased in the dataset. Also, we analysed the convergence behaviour of the Hopfield network by plotting histograms of update step distributions across varying corruption levels.




Fig 7. Required Update steps against frequency for each testing probability for the 25 Images Dataset

Next, we demonstrate the sequential evolution of states via the Hopfield network through the following image on various scenarios possible for listed pattern corruption techniques and update mechanisms.




Fig 8. Synchronous Update on bit flipped corrupted image for 25 Images Dataset

The original image was corrupted by flipping the bits randomly with probability $p = 0.3$. The synchronous update required 7 steps to reach the convergence. However, the network failed to capture the aspects of the test image.




Fig 9. Asynchronous Update on bit flipped corrupted image for 25 Images Dataset

The original image was corrupted by flipping the bits randomly with probability $p = 0.3$. The asynchronous update required 4 steps to reach the convergence for the same image. The image was not recovered in this case either.



Fig 10. Synchronous Update on the cropped image with bounding box for 25 Images Dataset

The original image was cropped by a 10 X 10 bounding box, where we kept part of the image inside the bounding box to be identical to the original image while changing every pixel outside to black. The synchronous update rule was applied, which required 7 steps to reach the convergence, however, the image was not fully recovered in this case.




Fig 11. Asynchronous Update on the cropped image with bounding box for 25 Images Dataset

The original image was cropped by a 10 X 10 bounding box, where we kept part of the image inside the bounding box to be identical to the original image while changing every pixel outside to black. The asynchronous update rule was applied, which required 4 steps to reach the convergence for the same image, however, the image was not fully recovered in this case, which could be attributed to the loss of key distinguishable features present in the bounding box or the bias of the network towards a spurious state.

We infer that the Hopfield Network failed to accurately recall test pictures after training with 25 patterns as it exceeds the theoretical capacity for training images^{[2][3]} of 16 X 16 pixels using the Hebbian Learning Rule^[1], resulting in inaccurate attractor states. The increased number of stored patterns may have resulted in overlapping memory representations, which prevented proper reconstruction, and might have forced them to settle into false states rather than accurate memories.

After the poor performance of the Hopfield network on a dataset of 25 images, We reduced the dataset to 10 images, model performance improved significantly which is evident from the findings presented below, however, a sustainable model with the given dataset size of 10 images was still not the most precise and was susceptible to excessive corruption and noise.

As with the entire dataset of 25 images, we first captured the average number of convergence steps against the corruption probability for a dataset with 10 images.




Fig 10. Average convergence Steps vs. Corruption probability for 10 Images dataset

Subsequently, in tandem with the previous experimentation with the 25-image dataset, we chalked a fraction of correctly recovered images against their corruption probability.




Fig 11. Fraction of correctly recovered images vs. Corruption probability for 10 Images Dataset

The fraction of correctly recovered images was significantly higher than in the previous case with partial recovery even for very high corruption probabilities such as 0.9 showcasing the improvement in the model performance.

Simultaneously, we analyzed the convergence behaviour of the Hopfield network by plotting histograms illustrating the distribution of update steps and the recovery accuracy across various corruption levels, leading to the following findings.




Fig 12. Required Update steps against frequency for each testing probability for the 10 Images Dataset

In parallel with the previous case with 25 images, the image recovery updates in the Hopfield Network were printed. For this case, we were able to recover the images both with synchronous and asynchronous methods while corrupting the bits of the image, however, in the cropping test, the images could not be recovered. A sample test image to support the claim has been shown below.




Fig 13. Synchronous Update on bit flipped corrupted image for 10 Images Dataset

The original image was distorted by randomly flipping the bits with a frequency of 0.3. The synchronous update needed 3 steps to achieve successful convergence.




Fig 14. Asynchronous Update on bit flipped corrupted image for 10 Images Dataset

Similarly, the asynchronous update rule was applied to the corrupted image with the probability of corruption being 0.3, which also reached successful convergence in 2 steps.




Fig 15. Synchronous Update on the cropped image with bounding box for 10 Images Dataset

The original image was cropped using a 10 X 10 bounding box, with part of the image remaining identical to the original image while every pixel outside was changed to black. The synchronous update rule was used, which needed 4 steps to achieve convergence; nonetheless, we were not able to completely restore the image.




Fig 16. Asynchronous Update on the cropped image with bounding box for 10 Images Dataset

Similarly, the asynchronous update rule was applied, which also required 4 steps to reach convergence but could not fully recover the image.

A key observation was made that the number of convergence steps needed was reduced in comparison to the complete dataset of 25 images for the test image, additionally, we were able to recover the image for the synchronous update. However, further investigation was needed to showcase the amount of data a Hopfield network can handle after introducing corruptions and cropping.

Therefore, for further evaluation, we trained the built Hopfield Network on 5 images and we present our findings below.

We visualised average convergence steps reached across different image bit corruption probabilities to analyse the performance of the constructed network on both synchronous and asynchronous updates.

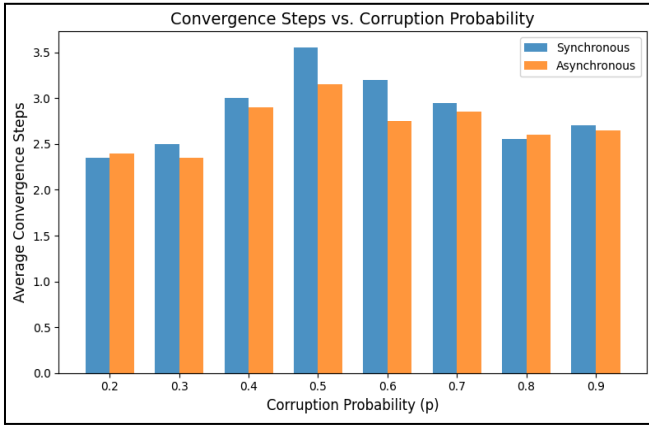


Fig 17. Average Convergence Steps vs. Corruption Probability for 5 images dataset

We find that the interference between stored memories might be possible since the network was trained on only 5 images, which is pronounced especially at moderate corruption levels, leading to slower convergence.

Moreover, we investigated the recall performance showcasing a fraction of successfully retrieved patterns across different image bit corruption rates for synchronous and asynchronous updates.




Fig 18. Fraction of correctly recovered images vs. Corruption probability for 5 images dataset

We also infer that the presence of structured patterns might help in easier recovery at low corruption levels but could also introduce spurious attractors at high corruption levels, potentially leading to incorrect convergence.

We also present the histograms of update step distributions and recovery accuracy across different levels of corruption to analyse the Hopfield network's convergence behaviour on the dataset of 5 images.




Fig 19. Required Update steps against frequency for each testing probability for the 5 Images Dataset

Next, we show a test image to illustrate the sequential evolution of states via the Hopfield network for the dataset of 5 images across the aforementioned corruption approaches and update processes.




Fig 20. Synchronous Update on bit flipped corrupted image for 5 Images Dataset

Similar to previous cases, the original image was distorted by randomly flipping the bits with a frequency of 0.3. The synchronous update took 3 steps to achieve convergence.




Fig 21. Asynchronous Update on bit flipped corrupted image for 5 Images Dataset

Similarly, random bit flipping with probability $p = 0.3$ was used to distort the original picture, for which the asynchronous update required 3 steps for convergence.




Fig 22. Synchronous Update on the cropped image with bounding box for 5 Images Dataset

In tandem with the previous cases, a 10×10 bounding box was used to crop the original image where we kept part of the image inside the bounding box to be identical to the original image while changing every pixel outside to black. It took 3 steps to obtain convergence when we used the synchronous update rule.




Fig 23. Asynchronous Update on the cropped image with bounding box for 5 Images Dataset

Similarly, the original image was cropped using a 10 X 10 bounding box, with part of the image remaining identical to the original image while every pixel outside was changed to black. The asynchronous update rule was used, which took 3 steps to achieve convergence for the same picture.

After analysing the various cases described above, we conclude that the Hopfield Network trained on a reduced image dataset consistently achieved successful convergence across all cases, which highlights that a limited number of stored patterns enables effective recall even when the images are moderately corrupted. However, as corruption increased, recovery rates dropped sharply, likely due to spurious attractors leading to incorrect convergence. Despite these challenges, the network consistently reached convergence, indicating that reducing the stored images provides a practical trade-off between storage capacity and recall reliability in the built Hopfield model for a dataset consisting of 16 X 16 pixels images.

IV. REFERENCES

1. Hebb, D. O. The organization of behavior: A neuropsychological theory. New York: John Wiley and Sons, Inc., 1949. 335 p. \$4.00. (1950). *Science Education*, 34(5), 336–337. <https://doi.org/10.1002/sce.37303405110>
2. Storkey, A. (1997). Increasing the capacity of a hopfield network without sacrificing functionality. In *Lecture notes in computer science* (pp. 451–456). <https://doi.org/10.1007/bfb0020196>
3. McEliece, R., Posner, E., Rodemich, E., & Venkatesh, S. (1987). The capacity of the Hopfield associative memory. *IEEE Transactions on Information Theory*, 33(4), 461–482. <https://doi.org/10.1109/tit.1987.1057328>

APPENDIX

The GitHub link along with relevant files for the source code in Python is provided below:

[bPratyush/CS-F433-Computational-Neuroscience-Assignments: Coursework Assignments for CS F433 Computational Neuroscience](https://github.com/Pratyush/CS-F433-Computational-Neuroscience-Assignments)

The raw code is provided below:

```
import numpy as np
import glob
import matplotlib.pyplot as plt
def readfile(filename):
    with open(filename,'r') as f:
        lines=f.readlines()
        lines=[l.strip() for l in lines if not l.startswith("#")]
        assert lines[0] == 'P1'
        width, height=map(int, lines[1].split())
        data=np.array([int(b) for line in lines[2:] for b in line.split()])
        return data.reshape((height, width))
def conv(mem):
    return 2*mem-1
# PROBLEM STATEMENT 1 Hopfield Network Training using Hebbian Rule
def trainingnet(memories):
    n=memories[0].size
    W=np.zeros((n, n))
    for mem in memories:
        vec=mem.flatten()
        W+= np.outer(vec, vec) #Hebbian Learning Rule
    W/=len(memories)
    np.fill_diagonal(W,0)
    return W
def bitflip(mem,p): #Bit flip corruption
    corrupt =np.random.rand(*mem.shape) < p
    corrupted_mem=np.copy(mem)
    corrupted_mem[corrupt] *= -1
    return corrupted_mem
def cropping(mem, bbox_color): #Bounding box cropping
    cropped = np.full(mem.shape, -1 if bbox_color == "black" else 1)
    cropped[3:13, 3:13] = mem[3:13, 3:13]
    return cropped
def syncupdate(state,W): #Synchronous Update
    h, w=state.shape
    flat=state.flatten()
    updated= np.sign(W @ flat)
    updated[updated==0] = 1
    return updated.reshape(h, w)
def asyncupdate(state,W): #Asynchronous Update
    h, w=state.shape
    flat=state.flatten()
    indices=np.arange(flat.size)
    np.random.shuffle(indices) #Random Shuffle
    for i in indices:
        flat[i]=1 if np.dot(W[i],flat) >= 0 else -1
    return flat.reshape(h, w)
def runconv(initial, W,updating): #Run until we reach convergence
```

```

prev_state=None
state=initial
steps=0
hist =[initial.copy()]
while not np.array_equal(state,prev_state):
    prev_state=state.copy()
    state =updating(state, W)
    steps+= 1
    hist.append(state.copy())
return state, steps, hist
# UTILITY FUNCTIONS FOR PLOTTING THE CURVES
def plotseq(hist,title): #Sequence Plotting
    n=len(hist)
    fig, axes=plt.subplots(1,n,figsize=(n * 3, 3),squeeze=False)
    for i, state in enumerate(hist):
        ax = axes[0, i]
        ax.imshow(state,cmap='gray',interpolation='nearest')
        ax.set_xticks([])
        ax.set_yticks([])
        ax.set_title(f"Step {i}", fontsize=10)
        for spine in ax.spines.values():
            spine.set_visible(False)
    fig.suptitle(title,fontsize=14)
    plt.tight_layout(rect=[0, 0, 1, 0.93])
    plt.show()
def plothist(stepcnt,pval): #Histogram Plotting
    fig,axes=plt.subplots(1,len(pval),figsize=(15, 4),sharey=True)
    for ax, p in zip(axes,pval):
        ax.hist(stepcnt[p],bins=6,alpha=0.5,edgecolor='black',color='skyblue')
        ax.set_title(f'p = {p}')
        ax.set_xlabel("Update Steps")
        ax.set_ylabel("Frequency")
    plt.tight_layout()
    plt.show()
def fillmemory(files):
    memories=[conv(readfile(f)) for f in files[:10]] #Change [:10] to [:5] for 5 images and remove slicing for 25 images
    return memories
def plotconvsteps(pval,syncsteps,asyncsteps): #Convergence Steps Plotting
    plt.figure(figsize=(8,5))
    w=0.35
    indices = np.arange(len(pval))
    plt.bar(indices,syncsteps,w,label='Synchronous',alpha=0.8)
    plt.bar(indices + w,asyncsteps,w,label='Asynchronous',alpha=0.8)
    plt.xlabel("Corruption Probability (p)",fontsize=12)
    plt.ylabel("Average Convergence Steps",fontsize=12)
    plt.xticks(indices + w/2,[f"{p:.1f}" for p in pval],fontsize=10)
    plt.legend(fontsize=10)
    plt.title("Convergence Steps vs. Corruption Probability",fontsize=14)
    plt.tight_layout()
    plt.show()
def plotcorrect(pval,corrsync,corrasync): #Correctly Recovered Plotting
    plt.figure(figsize=(8,5))
    w=0.35
    indices = np.arange(len(pval))

```



```

plt.bar(indices, corrsync,w,label='Synchronous',alpha=0.8)
plt.bar(indices + w, corrasync,w,label='Asynchronous',alpha=0.8)
plt.xlabel("Corruption Probability (p)",fontsize=12)
plt.ylabel("Fraction Correctly Recovered",fontsize=12)
plt.xticks(indices + w/2,[f'{p:.1f}' for p in pval],fontsize=10)
plt.legend(fontsize=10)
plt.title("Fraction Correctly Recovered vs. Corruption Probability",fontsize=14)
plt.tight_layout()
plt.show()
# Saving the state of the Hopfield Network onto a PBM file midway during the updates as in PROBLEM STATEMENT 1
def savepbm(state,filename):
    # Convert from -1/1 to 0/1 for image reconstruction
    state_01=((state+1)/2).astype(int)
    h, w=state_01.shape
    with open(filename,'w') as f:
        f.write('P1\n')
        f.write(f'{w} {h}\n')
        for row in state_01:
            f.write(' '.join(str(val) for val in row)+ '\n')
# DRIVER CODE
def main():
    files=sorted(glob.glob('pattern_*.pbm'))
    memories=fillmemory(files)
    W=trainingnet(memories)

    # PROBLEM STATEMENT 2
    p=0.3
    memidx = np.random.randint(len(memories))
    orig=memories[memidx]
    bitflippy=bitflip(orig, p)
    croppy=cropping(orig, "black")

    finsyncflippy, syncstepflippy, histsyncflip = runconv(bitflippy, W, syncupdate) #Bit flipped with synchronous update
    finasyncflippy, asyncstepflippy, histasyncflip = runconv(bitflippy, W, asyncupdate) #Bit flipped with asynchronous update
    finsynccroppy, syncstepcroppy, histsynccrop = runconv(croppy, W, syncupdate) #Cropped with synchronous update
    finasynccroppy, asyncstepcroppy, histasynccrop = runconv(croppy, W, asyncupdate) #Cropped with asynchronous update

    print(f'Flipped corruption (p={p}): Synchronous steps = {syncstepflippy}, Asynchronous steps = {asyncstepflippy}')
    print(f'Cropped corruption: Synchronous steps = {syncstepcroppy}, Asynchronous steps = {asyncstepcroppy}')

    #Visualisations
    plotseq(histsyncflip, "Synchronous Update (Flipped Corruption)")
    plotseq(histasyncflip, "Asynchronous Update (Flipped Corruption)")
    plotseq(histsynccrop, "Synchronous Update (Cropped Corruption)")
    plotseq(histasynccrop, "Asynchronous Update (Cropped Corruption)")

    #Saving Midway updates to PBM files
    if histsyncflip:
        midpoint = len(histsyncflip)//2
        savepbm(histsyncflip[midpoint], 'midway_sync_flip.pbm')
        print("Saved midway synchronous flip state to 'midway_sync_flip.pbm'")
    if histasyncflip:
        midpoint = len(histasyncflip)//2
        savepbm(histasyncflip[midpoint], 'midway_async_flip.pbm')

```

```

    print("Saved midway asynchronous flip state to 'midway_async_flip.pbm'")
if histsynccrop:
    midpoint = len(histsynccrop)//2
    savepbm(histsynccrop[midpoint], 'midway_sync_crop.pbm')
    print("Saved midway synchronous crop state to 'midway_sync_crop.pbm'")
if histasynccrop:
    midpoint = len(histasynccrop)//2
    savepbm(histasynccrop[midpoint], 'midway_async_crop.pbm')
    print("Saved midway asynchronous crop state to 'midway_async_crop.pbm'")
# PROBLEM STATEMENT 3
pval=[0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9]
paircntp=20
syncstepcnt={p: [] for p in pval}
asyncstepcnt={p: [] for p in pval}
corrsync={p: [] for p in pval}
corrasync={p: [] for p in pval}
for p in pval:
    for _ in range(paircntp):
        memidx=np.random.randint(len(memories))
        orig=memories[memidx]
        corrupted=bitflip(orig, p)
        finsync, stepsync, _ = runconv(corrupted, W, syncupdate)
        finasync, stepasync, _ = runconv(corrupted, W, asyncupdate)
        syncstepcnt[p].append(stepsync)
        asyncstepcnt[p].append(stepasync)
        corrsync[p].append(np.mean(finsync==orig))
        corrasync[p].append(np.mean(finasync==orig))
avgsyncsteps=[np.mean(syncstepcnt[p]) for p in pval]
avgasyncsteps=[np.mean(asyncstepcnt[p]) for p in pval]
avgcorrsync=[np.mean(corrsync[p]) for p in pval]
avgcorrasync=[np.mean(corrasync[p]) for p in pval]
plotconvsteps(pval,avgsyncsteps,avgasyncsteps)
plotcorrect(pval,avgcorrsync,avgcorrasync)
plothist(syncstepcnt,pval)
if __name__=="__main__":
    main()

```